

Controlled ML Study

Abstract

We evaluate a compact classifier on two public datasets.

1 Method

The method adds a calibrated linear head to a frozen encoder.

2 Experiments

Table 1 reports mean accuracy over five fixed seeds.

Results and Limitations

Table 1: Baseline 71.0; proposed method 74.5.

Ablation: removing calibration lowers accuracy to 72.2.

Limitations

The study covers classification and two datasets only.